

Analysis of Attack on Tor

Andrew Rosen

November 18, 2014

This document provides a short summary and analysis of the work presented by Bauer *et al.* [1] on compromising the anonymity of Tor [2] users by controlling entry and exit nodes in the system.

1 Attack Description

The Low-Resource Routing Attack (LORA for brevity), is an attack that can to compromise the anonymity of a sender using the Tor network. In a typical Tor setup, a client creates a circuit through the Tor network to communicate with a destination. The goal is to prevent a nosy individual or organization from snooping on the message contents, as well as hiding who the user is communicating with. The client selects three routers¹ for the circuit: an entrance node, a mix node, and an exit node.

Messages are sent through the Tor network to the destination along the circuit, encrypted in multiple layers, such that each hop in Tor removes a layer of that encryption, until the contents are sent unencrypted² from the exit node to the destination.

A client selects the entry nodes by obtaining a small list of fast and stable routers from a trusted directory server. Here, fast and stable mean having more bandwidth and greater uptime than the median of the routers. The directory server obtains this information by listening to advertisements from routers.

Mix and exit nodes are also picked based on their bandwidth and stability, although the selection is random to avoid all circuits going through the same router. The greater the router's advertised bandwidth, the more likely a node is selected as a router.

Only the message sender knows the complete circuit the message took in the Tor network. Thus, an observer at the sender's side could tell the client was sending messages, but not what or to whom. An observer at the receiver's end would just see a bunch of random messages. Naturally this entire arrangement is thoroughly upsetting to snoopers.

2 The Attack

The issue with the above scheme is when an adversary controls both the entrance and exit of the network. The entry node knows who the sender is and the

¹Three is the default number; the original paper explored the default Tor settings.

²Well, not encrypted by the Tor network; the message itself may be some ciphertext for the destination to decrypt.

mix server the message is being passed on to. The exit node similarly knows which mix server the message is coming from and the final destination of the message. An adversary controlling both would be able to tell who a client is communicating with, although not necessarily what about, but that's enough to break anonymity.

The idea behind LORA is to game the router selection process so, with uncomfortable levels of probability, users select malicious nodes as the entrance and exit nodes of a Tor circuit. As discussed, routers along the circuit in Tor are selected based on their speed and stability.

The issue is that information about each router is provided by that router. This information is trusted but unverified, which means an attacker can advertise false speeds and uptime, a technique also known as *lying*.

The adversary has multiple routes for launching this attack. If he has the hardware, he can deploy high performing routers on the network, which would be more likely to be chosen, but probably not likely enough. If the routers don't have the capability needed to have a high chance of being chosen, the malicious routers can lie and say they can handle a disproportionate high bandwidth to skew the exit node selection in his favor.

The adversary can also over-represent his routers in the network using a Sybil [3] attack. In a Sybil attack, a user lies and says that they are multiple nodes. For example, a malicious router could join the network multiple times along the same IP but different ports. A well-executed Sybil attack can turn a non-global adversary into a global-adversary.

Finally, if only one malicious node is chosen as part of the circuit, the adversary can force the user to attempt select a new circuit by deliberately dropping the connection. Experiments demonstrated that an adversary controlling 10% of the routers in this manner could compromise about a third of the circuits formed.

3 Proposed Defenses

The defenses focused on building a reputation system that will eventually exclude routers falsely advertising their bandwidth. A client keeps a local reputation score for each router it can observe. The router's reputation score starts at its chances of being chosen as router for the circuit.

Whenever a circuit performs at or better than the minimum advertised bandwidth for the circuit, all the participating routers are rewarded with an increase to their reputation. Otherwise, all the circuit's routers' reputations are penalized. Routers with a low enough reputation are blacklisted.

The adjustment to the reputation is weighted such that greater the node's advertised bandwidth, the greater the adjustment. This means that routers participated in LORA are quickly blacklisted for their wildly inaccurate advertisements. To speed convergence to a blacklist, clients can pseudonymously share the observed reputations between each other and aggregate the results.

The reputation system can greatly impede the efforts of an adversary using LORA to attack client anonymity, but only LORA. It is still possible to perform this traffic attack with high resource machines, albeit more expensive.

For example, if adversary has compromised a circuit's entry and exit, so long as they perform better than or at the mix router's level, they will gain

reputation. In theory, an attacker could Sybil a high performing machine into place and only accept a circuit that has been compromised (since there is no penalty for dropping the circuit) and has a mix router of a desired bandwidth. To maximize the number of circuits, each circuit the malicious router is participating in could probe the mix router and limit the entry and exit's bandwidths to that of the mix routers.

This defense does not mitigate the threat of a Sybil attack, which would need some policy changes. Bauer *et al.* propose putting some filters on the directory servers, so that an IP or subnet can only host a single router.

This is in line with research done in making Sybil attacks more expensive for distributed hash tables (DHTs) [3] by limiting the number of repeated IPs in the network. One proposed defence for DHTs is to require joining nodes to register with some members of the network [4]. If the registration reveals there are too many nodes with that IP, it is rejected.

3.1 Other Defenses

Snader *et al.* proposes [5] a technique replacing self-advertisement with a distributed opportunistic measurement by other nodes. This was initially rejected by Bauer *et al.* due to the additional load on the network, but if probes are opportunistic, the bandwidth impact is minimal. This is combined with a new selection algorithm that selects routers in the circuit based on their ranking according to metric, which impedes an adversary hoping to game the system by using a high-bandwidth machine. However, this technique seems to still be vulnerable to a Sybil attack with a high bandwidth machine, since an attacker would get multiple entries in the ranking system.

4 Conclusion

The LORA attack is genuine threat to the anonymity of Tor clients, as it makes traffic interception at the entry and exit routers easy. Using a reputation system helps defend against the LORA attack specifically, but traffic interception is still possible. A major threat against Tor is the Sybil attack and the easiest way to make that impractical is to filter additions to the directory servers so that only unique IPs or subnets can be routers.

References

- [1] K. Bauer, D. McCoy, D. Grunwald, T. Kohno, and D. Sicker, "Low-resource routing attacks against anonymous systems,"
- [2] R. Dingledine, N. Mathewson, and P. Syverson, "Tor: The second-generation onion router," tech. rep., DTIC Document, 2004.
- [3] G. Urdaneta, G. Pierre, and M. V. Steen, "A survey of dht security techniques," *ACM Computing Surveys (CSUR)*, vol. 43, no. 2, p. 8, 2011.
- [4] J. Dinger and H. Hartenstein, "Defending the sybil attack in p2p networks: Taxonomy, challenges, and a proposal for self-registration," in *Availability*,

Reliability and Security, 2006. ARES 2006. The First International Conference on, pp. 8–pp, IEEE, 2006.

- [5] R. Snader and N. Borisov, “A tune-up for tor: Improving security and performance in the tor network.,” in *NDSS*, vol. 8, p. 127, 2008.