

The Sybil Attack on Peer-to-Peer Networks From the Attacker’s Perspective

Andrew Rosen

December 3, 2014

Abstract

This paper explores the feasibility of performing naive Sybil attacks on a P2P network that completely occlude healthy nodes from each other. The vulnerability of Distributed Hash Tables to Sybil attacks and Eclipse attacks have been well known for some time. However, these vulnerabilities has often been explored in a theoretical sense, assuming the attacker is an all-knowing and all-powerful global adversary from the beginning. This paper proves that these are not necessary characteristics and assumptions for an adversary to have to perform a Sybil attack.

We examine the amount of computational effort required to perform an Eclipse attack using only Sybil identities and without poisoning routing tables via false advertisement. We do this by analyzing the amount of effort it takes for an attacker with a given IP to choose a port to obtain a desired hash key, a process we call *mashing*. Mashing requires no additional resources on the adversary’s end. While mashing can be used for malicious behavior, we describe methods and applications for the technique that are beneficial for performing load balancing. This provides the flexibility for any node to effortlessly shoulder additional work from bottleneck locations if acting for the good of the network.

1 Introduction

One of the key properties of structured peer-to-peer (P2P) systems is the lack of a centralized coordinator. P2P systems remove the vulnerability of a single point of failure and the susceptibility to a denial of service attack (cite original Sybil paper), but in doing so, open themselves up to new attacks. Completely decentralized P2P systems are vulnerable to *Eclipse attacks*, whereby an attacker completely occlude healthy nodes from one another. This prevents them from communicating without being intercepted by the attacker.

One way to accomplish this attack is to perform a *Sybil Attack* [?]. In a Sybil attack, the attacker creates multiple malicious virtual nodes in order to disrupt the network. If enough malicious nodes are injected into the system, the majority of the nodes will be occluded from one another, successfully performing an Eclipse attack.

Security analyses typically assumes an adversary using a Sybil is omniscient and can inject virtual nodes wherever he chooses in a reasonable amount of time. Our goal is to demonstrate how veracity of that assumption by doing analysis and simulations.

Sybil attacks represent a significant threat to the security of any distributed system. Many of the analyses [1] on Tor [3] emphasize the vulnerability of Tor to the Sybil attack. This threatens the anonymity of Tor users, particularly those living in countries with peculiar notions about personal privacy.

P2P systems like BitTorrent are essential to a wide variety of users. BitTorrent, for instance, is the *de facto* platform for distributing large files scalably among tens of thousands of clients. An estimated 20 million peers use BitTorrent daily for sharing and retrieving Linux and BSD-based distros, offline copies of Wikipedia, and personal backup. Current research demonstrates BitTorrent is vulnerable to the Sybil attack and a persistent attack disabling BitTorrent would be highly detrimental to many users, but especially developers and system administrators¹

There have been many suggestions on how to defend against Sybil attack, but there is no agreed upon “silver bullet” among researchers that should be implemented for every distributed application [6]. Part of this is surely influenced by the only surefire way to defend against a Sybil attack is to introduce a trusted authority to certify and/or bind identities. This solution potentially removes the Sybil attack, but reintroduces vulnerabilities to denial of service attacks, bringing us full circle. Another reason is that there is no single solution for all platforms is that not every solution is compatible with each distributed platform.

Despite the threat represented by the Sybil attack and the research done on the subject, little research has been done from the perspective of an adversary. We sought to rectify this, both to reemphasize the threat of the Sybil attack, but also because this examination introduces some interesting graph theory problems.

Our work presents the following contributions:

¹Perhaps the only reason that BitTorrent hasn’t be attacked in such a way as to render it unusable is that those capable of doing so rely heavily upon it. A sobering thought, since BitTorrent is under an active Sybil attack [?].

- We first discuss the mechanics of performing a Sybil attack and analyze the theoretical effort needed to bring to bear to perform the Sybil attack.
- We present our simulations which show how quickly even a naive and inefficient Sybil attack can compromise a system. We also discuss how a more intelligent attack can be geared to each of the more popular DHT topologies.
- We discuss the implications of our work, specifically how the techniques we developed to perform the attack can be used for automatic load balancing.

2 Formal Analysis

2.1 Assumptions

In our analysis we make a few assumptions. First, we look primarily at fully distributed systems that assign nodes identifiers using a cryptographic hash function. These are called distributed hash tables (DHTs). Cryptographic hash functions work by mapping some input value to an m -bit key or identifier. Well-known hash functions for performing this task include MD5 [7] and SHA1 [5]. Keys generated by the hash function are assumed to be evenly distributed and random [2].

In distributed hash tables, m is a large value, in order to avoid collisions between mapped inputs, unintentional or otherwise. An $m \geq 128$ is typical, with $m = 160$ being the most popular choice.

Our second assumption is that node IDs are generated by hashing their IP address and port. The majority of DHTs present this method as the means of generating node IDs, and although other methods exist,² they are not nearly as prevalent.

We also assume that nodes store verify that a node's advertised IP address and port generate the advertised node ID. Relaxing the latter assumptions would make the Sybil attack trivial. If the attacker does not have to search for values, he can make up any needed key on the fly and falsely advertise it to the network unchallenged.

Consider an attacker operating under these assumptions who wants to inject a malicious node in between two victims. The attacker must search for a valid IP and port combination under his control that generates a hash key that lies between the two victims's keys, a process we will call *mashing*. The attackers

²In Mainline DHT, used by BitTorrent, it appears you can choose your own ID at "random." The horrifying ramifications should be immediately apparent and are discussed in Section 4.

ability to compromise the network depends on what IP addresses and ports he has available.

2.2 Analysis

Suppose we have a DHT with n members in it, with m -bit node IDs between $[0, 2^m)$. Consider two victim nodes with IDs a and b . The probability P that an attacker can mash a hash key that lands in the range (a, b) is

$$P \approx \frac{|b - a|}{2^m} \cdot \text{num_ips} \cdot \text{num_ports}$$

Where num_ip addresses is the number of IP addresses the attacker has under his control and num_ports is the number of ports the attacker can try for each IP address. If the ports the attacker can try are limited the ephemeral ports, the attacker has 16383 ports to use for each IP address. We can assume that, for a large enough n , node IDs will be close to evenly distributed across the keyspace, meaning there will be $\approx \frac{2^m}{n}$ key between each node ID. This makes the earlier probability equivalent to:

$$P \approx \frac{1}{n} \cdot \text{num_ips} \cdot \text{num_ports}$$

This indicates that the ease of mashing is independent to m .

The chances of an attacker mashing all n nodes that partition the entire keyspace is:

$$C \approx 1 - \left(1 - \frac{1}{n}\right)^{\text{num_ips} \cdot \text{num_ports}}$$

3 Simulations

An essential part of our analysis was demonstrating just how fast an adversary can compromise a system. We performed three experiments to accomplish this task. Simulations were performed in Python 2.7 with an AMD Phenom 965 processor.

We used SHA1 for as our cryptographic function, which yields 160-bit hash keys. Again, we use the constraint that victim nodes do an absolute minimum verification which forces an attacker to only present Sybils with hash keys that he can generate with valid IP and port combinations.

Another important aspect of our simulation is that we do not include the memory resources needed in this analysis. This is for a couple of reasons, the primary one being that we assume the attacker is capable and competent. An adversary who wants to use a Sybil attack on a large system would likely roll

his own client capable of running on the P2P protocol. Previous research [?] shows that an adversary performing a Sybil attack does not need to maintain the state information of other nodes and can leverage the healthy nodes to perform the routing. Nor does the adversary necessarily create a new running copy for every Sybil created.

If this was impossible, it is still reasonable to assume the attacker’s program would be built to have a minimal memory footprint and drop any messages that require an attacker carry any of the network’s load.

3.1 Experiment 1: mashing 2 random nodes

Our initial experiment was designed to establish the feasibility of injecting in between two random nodes. Each trial, we generated two victims with random IP addresses and ports, and an attacker with a random IP. The experiment was for the attacker to find a key in between the two victims’s key, from the lowest to highest. The amount of time to mash two random keys was on average 29.6218323708 microseconds, and was achievable 99.996% of the time.

3.2 Experiment 2: Nearest Neighbor Eclipse via Sybil

The objective of the second experiment is to completely ensnare a network using a Sybil attack, starting with a single malicious node. With simulate this by creating a network of n nodes. Each node is represented by a key generated by taking the SHA1 of a random IP/port combination.

The attacker is given a randomly generated *numIP* IP addresses, but can use any port between 49152 and 65535. This gives the attacker $16383 \cdot \text{numIPs}$ possible hash keys to use for Sybils. The attacker can easily precompute all of theses hash keys and store them in a sorted list to be used as needed. This requires only $160 \cdot 16383 = 2621280$ bits, or about 320 kilobytes for each set of hash keys generated for an IP.³

To perform the attack, an adversary choses any random hash key as a starting point to “join” the network. This is his first Sybil and the join process provides information about a number of other nodes. Most importantly, nodes provide information about other nodes that are close in the hash space, which provided to enable fault tolerance between immediate neighbors. The adversary uses this information to inject Sybils in between successive healthy nodes.

For clarity we present this example. Consider the small segment of the

³We can ignore the cost of having to store 32 bits for each IPv4 address and 16 bits required by each port with some clever mapping and indexing.

network made up of adjacent nodes a , b , c , and d . The Sybil joins between nodes a and b , and the joining process informs the adversary about node c , possibly node d , and a handful of other nodes in the network. The adversary will always learn about node c because a normal node between a and b would need to know about node c for fault tolerance.

The adversary’s next move would be to inject a node between nodes b and c . This is done by selecting a precomputed hash key k , such that $b \leq k \leq c$. The adversary injects a Sybil node with key k , which joins in between b and c , and the joining process informs the adversary about node d and several other nodes, including many close nodes. The adversary then aims to inject a node in between c and d , and continues ad nauseum.

Depending on the network size and the number of keys available to the adversary, it is entirely possible the adversary will not have a key to inject between a pair of successive nodes. In this case, the node moves on to the successive pair that the node has learned about.

We simulated this attack on networks of size 50, 100, 200, 300, 400, 500, 1000, 5000, 10000, 15000, 20000, 25000, 50000, and 100000. We gave the attack access to 1, 3, 5, and 10 IP addresses. We presents our results in Table 3.2; we have included only the results for larger sized networks, as the smaller sized networks were completely dominated.

Sybil IPs	Network Size	% of Regions Injected	Sybils/Region
1	25000	0.39556	0.65532
1	50000	0.24638	0.32766
1	100000	0.14065	0.16383
3	25000	0.66068	1.96592
3	50000	0.49712	0.98294
3	100000	0.32942	0.49148
5	25000	0.76964	3.2766
5	50000	0.62334	1.63824
5	100000	0.45032	0.81912
10	25000	0.86792	6.55292
10	50000	0.76584	3.27648
10	100000	0.62321	1.63829

Table 1: Results for Sybil Attack. The Fraction of regions injected is the percentage of all regions for which the adversary could inject a Sybil with the suitable key. The Sybil/Region measure is how many Sybils were available to inject a particular region on average.

Our results show that an adversary, given only modest resources, can inject a Sybil in between the vast majority of successive nodes in moderately sized networks. In a large network, modest resources still can be used to compromise more than a third of the network, an important goal if the adversary wishes to launch a Byzantine attack. The Sybil attack degraded below that point for

However, this experiment only covers the short links of the network, but not the long distance link.

Given uncompromised node, what is the chances the node is a Sybil

$$\frac{SybilKeys}{SybilKeys + NetworkSize - 1}$$

3.3 Experiment 3: Fully Complete Eclipse via Sybil

We extended the previous experiment by considering the long-distance hops of each node in addition to the short-range links of the DHT. We choose to model an attack on a P2P system built using Chord [?]. We did this for a number of reasons. Chord is an extremely well understood DHT and is very simple to evaluate using simulations. Nodes in Chord generate long distance links independent of information provided by other nodes, rather than directly querying neighbors. This prevents adversaries from poisoning the node's routing table via false advertisements, which can be on other DHTs such as Pastry [?]. A Sybil attack the most straightforward means of effecting an Eclipse attack on a Chord network.

Nodes in Chord have m long-range links, one for each of the bits in the keys, which is 160 in our experiments. Each of node a 's long-range hops points to the node with the lowest key $\geq a + 2^i \bmod 2^m, 0 \leq i \leq 160$.

The attack is very similar to the Nearest Neighbor attack demonstrated above. Beside injecting a node in between successive nodes, the attacker also attempts to place a Sybil in between each of the long-range links. We simulated this attack under the same parameters as above and are presented in Table 3.3.

4 Conclusions and Future Work

Our analysis and experiments show that an adversary with a small number of machines and limiting their port usage to just the ephemeral ports can easily compromise a P2P system and mash the majority of the regions between nodes. This effectively prevents nodes from talking to one another without first sending messages through Sybils, who can eavesdrop on them or disregard the eaves and move straight on to the dropping.

Sybil IPs	Network Size	% links occluded	Occlusion per node
1	50	1.0	160.0
1	500	0.9719375	155.51
1	1000	0.939525	150.324
1	5000	0.76638875	122.6222
3	50	0.99975	159.96
3	500	0.9899	158.384
3	1000	0.97068125	155.309
3	5000	0.9110975	145.7756
5	50	0.999875	159.98
5	1000	0.9899625	158.394
5	5000	0.9444825	151.1172
10	50	1.0	160.0
10	500	0.9960125	159.362
10	1000	0.99673125	159.477
10	5000	0.9693825	155.1012

Table 2: Results for Sybil Attack on Chord. The percentage of links occluded is the percentage of long range links from healthy nodes that connect to Sybil nodes. The Occlusion per node measures how many of the 160 long range links lead to a Sybil on average. We ignore the predecessor links since it is not used for routing.

Our discussion have primarily concerned Chord, but an astute reader may wonder why did we not simulate an attack on Mainline DHT (MLDHT) [?], the Kademlia [?] based DHT used as the backend of BitTorrent. It is ostensibly the largest P2P system in use, so it seems more suited for analyzing our mash-based attack than Chord.

We avoided simulating the mashing process on MLDHT because it is completely unnecessary to perform the attack. In MLDHT, a node ID is not chosen by hashing an IP and port combination, but by picking an address uniformly at random between 0 and $2^{160} - 1$. An adversary who wants to a completely isolate a node on the network can choose Sybils with “random” hashkeys completely eliminate that node from the routing tables of healthy peers. Research has examined Mainline’s vulnerability to Sybil attacks[?] and detected entities performing the attack on Mainline. Our research demonstrates that switching to using a SHA1 hash of the IP and port is no defense against a Sybil attack, since the adversary can easily mash into appropriate places.

However, the mashing process can be used in non-security related settings

to benefit a DHT. As we have mention previously, the SHA1 hash has an evenly distributed output, but no set of keys will be completely evenly distributed. Some nodes will be responsible for larger regions than others⁴ and be responsible for a larger portion for the work.

If a node can detect a when a peer has too much of a load, the node can inject a virtual node into the region to shoulder some of the load. Too much of a load could be defined by having a too large region or by watching traffic requests. This could be integrated into a distributed key allocation system, so that a node can bind these keys to itself, but only use a small portion of the keys allotted to it to aid with load-balancing. The only cost would be that node would have to keep track of the keys it can inject, which we mentioned was only 320 kilobytes of data.

References

- [1] Kevin Bauer, Damon McCoy, Dirk Grunwald, Tadayoshi Kohno, and Douglas Sicker. Low-resource routing attacks against tor. In *Proceedings of the 2007 ACM workshop on Privacy in electronic society*, pages 11–20. ACM, 2007.
- [2] Mihir Bellare and Tadayoshi Kohno. Hash function balance and its impact on birthday attacks. In *Advances in Cryptology-Eurocrypt 2004*, pages 401–418. Springer, 2004.
- [3] Roger Dingledine, Nick Mathewson, and Paul Syverson. Tor: The second-generation onion router. Technical report, DTIC Document, 2004.
- [4] John R Douceur. The sybil attack. In *Peer-to-peer Systems*, pages 251–260. Springer, 2002.
- [5] Donald Eastlake and Paul Jones. Us secure hash algorithm 1 (sha1), 2001.
- [6] Brian Neil Levine, Clay Shields, and N Boris Margolin. A survey of solutions to the sybil attack. *University of Massachusetts Amherst, Amherst, MA*, 2006.
- [7] Ronald Rivest. The md5 message-digest algorithm. 1992.

⁴More susceptible to a mash.